



Министерство образования и науки Российской Федерации
Федеральное государственное бюджетное образовательное
учреждение высшего образования
«Московский государственный технический университет
имени Н.Э. Баумана
(национальный исследовательский университет)»
(МГТУ им. Н.Э. Баумана)

МЕТОДЫ МАШИННОГО ОБУЧЕНИЯ

Отчёт по рубежному контролю № 2

*«Задача классификация текстов
на основе методов наивного Байеса»*

Выполнила:

студентка группы ИУ5 – 23М

Галичий Д. А.

Преподаватель:

Гапанюк Ю. Е.

2020г.

```

import numpy as np
import pandas as pd
from typing import Dict, Tuple
from scipy import stats
from IPython.display import Image
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsRegressor, KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.metrics import accuracy_score, balanced_accuracy_score
from sklearn.metrics import precision_score, recall_score, f1_score, classification_report
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import cross_val_score
from sklearn.pipeline import Pipeline
from sklearn.metrics import mean_absolute_error, mean_squared_error, mean_squared_log_error
from sklearn.metrics import roc_curve, roc_auc_score, auc
from sklearn.svm import SVC, NuSVC, LinearSVC, OneClassSVM, SVR, NuSVR, LinearSVR
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(style="ticks")

```

```

# В качестве набора данных будем использовать отзывы с сайта amazon
mob = pd.read_csv('Amazon_Unlocked_Mobile.csv', sep=",", encoding='ISO-8859-1')
mob.head()

```

		Product Name	Brand Name	Price	Rating	
0	"CLEAR CLEAN ESN" Sprint EPIC 4G Galaxy SPH-D7...	Samsung	199.99	5	I feel so	
1	"CLEAR CLEAN ESN" Sprint EPIC 4G Galaxy SPH-D7...	Samsung	199.99	4	nice pho	
2	"CLEAR CLEAN ESN" Sprint EPIC 4G Galaxy SPH-D7...	Samsung	199.99	5		
3	"CLEAR CLEAN ESN" Sprint EPIC 4G Galaxy SPH-D7...	Samsung	199.99	4	It wor	
4	"CLEAR CLEAN ESN" Sprint EPIC 4G Galaxy SPH-D7...	Samsung	199.99	4	Great p	

```

mob.shape

```

```

[4] (413840, 6)

```

PromptCloud extracted more than 400 thousand reviews of unlocked mobile phones sold on Amazon. Let's take a look at the reviews, ratings, price and their relationships.

```

# Уникальные значения признака Review Votes
mob['Review Votes'].unique()

```

```

[4]

```

```
array([ 1.,  0.,  2.,  6., 19., 12., 13., 17.,  4.,  9.,  3.,
       5., 16.,  7., 11., 15.,  8., 23., 14., 20., 24., nan,
      18., 10., 41., 59., 42., 116., 22., 57., 21., 32., 54.,
     30., 77., 115., 39., 31., 64., 124., 36., 100., 125., 28.,
     86., 25., 43., 27., 120., 72., 121., 48., 33., 75., 40.,
    37., 249., 38., 130., 112., 157., 51., 58., 63., 107., 52.,
    34., 46., 45., 85., 105., 366., 29., 138., 67., 62., 35.,
    47., 177., 80., 66., 68., 84., 152., 53., 113., 151., 141.,
   159., 128., 94., 79., 56., 78., 55., 61., 50., 76., 127.,
   96., 119., 87., 60., 192., 44., 49., 126., 26., 81., 88.,
  158., 92., 155., 154., 156., 106., 175., 103., 172., 118., 83.,
  296., 136., 73., 74., 90., 65., 283., 282., 288., 285., 270.,
  265., 478., 462., 465., 171., 98., 99., 169., 213., 110., 183.,
   71., 109., 114., 182., 221., 487., 170., 93., 166., 137., 108.,
  219., 123., 218., 91., 69., 185., 226., 122., 146., 190., 246.,
  250., 82., 150., 214., 133., 274., 101., 257., 129., 215., 104.,
  140., 134., 200., 297., 204., 176., 232., 227., 263., 102., 89.,
   70., 173., 153., 404., 339., 131., 524., 518., 331., 95., 519.,
  287., 187., 197., 117., 264., 186., 142., 193., 145., 147., 191.,
  207., 383., 251., 240., 139., 305., 304., 174., 132., 308., 97.,
  279., 216., 160., 178., 111., 428., 423., 205., 356., 645., 165.,
  168 . 212 . 371 . 295 . 364 . 163 . 201 . 272 . 162 . 350 . 335 1])
```

```
# Уникальные значения признака Rating
```

```
# Будем использовать этот признак для определения характеристик текста
```

```
mob['Rating'].unique()
```

```
↳ array([5, 4, 1, 2, 3])
```

```
mob.isna().sum()
```

```
↳ Product Name      0
  Brand Name      65171
  Price          5933
  Rating         0
  Reviews        62
  Review Votes  12296
  dtype: int64
```

```
# Удалим строки с пропусками данных,
```

```
# так как их количество существенно меньше размера датасета и не отразится на качестве мод
mob.dropna(inplace=True)
```

```
mob.shape
```

```
↳ (334335, 6)
```

Рассмотрим, как рейтинг мобильных устройств влияет на длину отзывов

```
s1 = mob[mob.Rating == 1].Reviews.apply(len)
s2 = mob[mob.Rating == 2].Reviews.apply(len)
s3 = mob[mob.Rating == 3].Reviews.apply(len)
s4 = mob[mob.Rating == 4].Reviews.apply(len)
s5 = mob[mob.Rating == 5].Reviews.apply(len)
```

```

sns.distplot(s1,label='1')
sns.distplot(s2,label='2')
sns.distplot(s3,label='3')
sns.distplot(s4,label='4')
sns.distplot(s5,label='5')

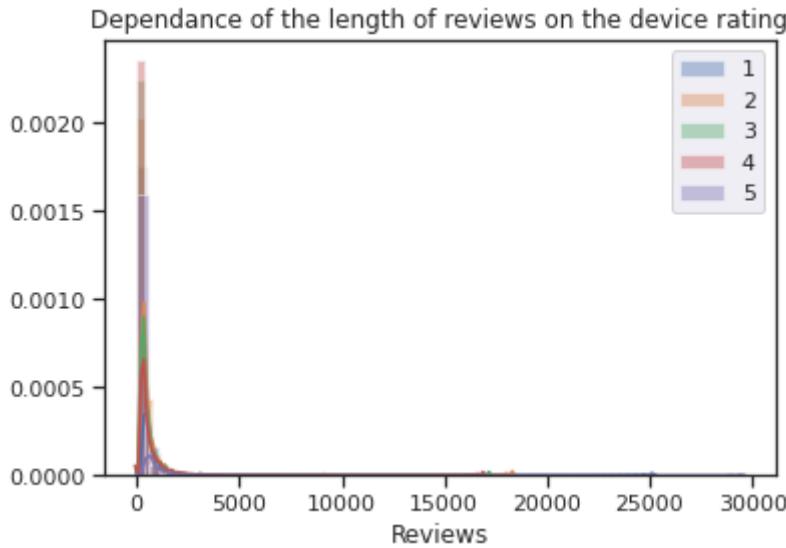
sns.set()

plt.title('Dependance of the length of reviews on the device rating')
plt.legend()

print('1 - mean: %s' % s1.mean())
print('2 - mean: %s' % s2.mean())
print('3 - mean: %s' % s3.mean())
print('4 - mean: %s' % s4.mean())
print('5 - mean: %s' % s5.mean())

```

→ 1 - mean: 247.22681845832972
 2 - mean: 330.6330476380307
 3 - mean: 296.54532197405786
 4 - mean: 274.9803058249539
 5 - mean: 169.16030246375928



Как видно из графика, негативные отзывы (для товаров с рейтингом 1,2) превышают положительными (рейтингом 4,5) по длине.

```

# Разделим данные на обучающую и тестовую выборки
X = mob.Reviews
y = mob.Rating
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)

```

```

# Сформируем общий словарь для обучения моделей из обучающей и тестовой выборки
vocab_list = mob['Reviews'].tolist()
vocab_list[1:10]

```

→

```
['nice phone, nice up grade from my pantach revue. Very clean set up and easy set up.  
'Very pleased',  
'It works good but it goes slow sometimes but its a very good phone I love it',  
'Great phone to replace my lost phone. The only thing is the volume up button does n  
'I already had a phone with problems... I know it stated it was used, but dang, it d  
'The charging port was loose. I got that soldered in. Then needed a new battery as w  
"Phone looks good but wouldn't stay charged, had to buy new battery. Still couldn't  
'I originally was using the Samsung S2 Galaxy for Sprint and wanted to return back t  
vocabVect = CountVectorizer()  
vocabVect.fit(vocab_list)  
corpusVocab = vocabVect.vocabulary_  
print('Количество сформированных признаков - {}'.format(len(corpusVocab)))
```

↳ Количество сформированных признаков - 63516

```
for i in list(corpusVocab)[1:10]:  
    print('{}={}'.format(i, corpusVocab[i]))
```

↳ so=52056
lucky=34179
to=56818
have=27444
found=24315
this=56266
used=59467
phone=41769
us=59365

Проведём векторизацию текста

```
# Класс TfidfVectorizer вычисляет специфичность текста в корпусе текстов на основе метрики  
tfidfv = TfidfVectorizer(ngram_range=(1,3))  
tfidf_ngram_features = tfidfv.fit_transform(vocab_list)  
tfidf_ngram_features
```

↳ <334335x4204950 sparse matrix of type '<class 'numpy.float64'>'
with 33222615 stored elements in Compressed Sparse Row format>

В качестве классификаторов будем использовать Multinomial Naive Bayes (MNB) и LogisticR

```
clf1 = Pipeline([('tfidfv', tfidfv), ('MultinomNB', MultinomialNB()),])  
clf2 = Pipeline([('tfidfv', tfidfv), ('LogRegr', LogisticRegression()),])
```

Multinomial Naive Bayes (MNB)

```
%time  
clf1.fit(X_train, y_train)
```

↳

```
CPU times: user 2 µs, sys: 0 ns, total: 2 µs
Wall time: 44.1 µs
Pipeline(memory=None,
         steps=[('tfidfv',
                 TfidfVectorizer(analyzer='word', binary=False,
                                 decode_error='strict',
                                 dtype=<class 'numpy.float64'>,
                                 encoding='utf-8', input='content',
                                 lowercase=True, max_df=1.0, max_features=None,
                                 min_df=1, ngram_range=(1, 3), norm='l2',
                                 preprocessor=None, smooth_idf=True,
                                 stop_words=None, strip_accents=None,
                                 sublinear_tf=False,
                                 token_pattern='(\\b\\w+\\b',
                                 tokenizer=None, use_idf=True,
                                 vocabulary=None)),
              ('MultinomNB',
               MultinomialNB(alpha=1.0, class_prior=None, fit_prior=True))])
```

Метрики точности accuracy_score, precision_score, recall_score

```
print('accuracy_score (train):', accuracy_score(y_train, clf1.predict(X_train)))
print('accuracy_score (test):', accuracy_score(y_test, clf1.predict(X_test)))
```

```
accuracy_score (train): 0.6868831024551988  
accuracy_score (test): 0.6610901187425848
```

```
print('precision_score (train):', precision_score(y_train, clf1.predict(X_train), average='macro'))  
print('precision_score (test):', precision_score(y_test, clf1.predict(X_test), average='macro'))
```

```
precision_score (train): 0.6868831024551988  
precision_score (test): 0.6610901187425848
```

```
print('recall_score (train):', recall_score(y_train, clf1.predict(X_train), average='micro'))  
print('recall_score (test):', recall_score(y_test, clf1.predict(X_test), average='micro'))
```

```
recall_score (train): 0.6868831024551988  
recall_score (test): 0.6610901187425848
```

LogisticRegression

```
%time  
clf2.fit(X_train, y_train)
```

1

```
CPU times: user 2 µs, sys: 0 ns, total: 2 µs
Wall time: 5.48 µs
/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/_logistic.py:940: Converg
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
Pipeline(memory=None,
          steps=[('tfidfv',
                  TfidfVectorizer(analyzer='word', binary=False,
                                  decode_error='strict',
                                  dtype=<class 'numpy.float64'>,
                                  encoding='utf-8', input='content',
                                  lowercase=True, max_df=1.0, max_features=None,
                                  min_df=1, ngram_range=(1, 3), norm='l2',
                                  preprocessor=None, smooth_idf=True,
                                  stop_words=None, strip_accents=None,
                                  sublinear_tf=False,
                                  token_pattern='(?u)\\b\\w\\\\w+\\b',
                                  tokenizer=None, use_idf=True,
                                  vocabulary=None))]
```

Метрики точности accuracy_score, precision_score, recall_score

```
fit intercept=True, intercept_scaling=1,
print('accuracy_score (train):', accuracy_score(y_train, clf2.predict(X_train)))
print('accuracy_score (test):', accuracy_score(y_test, clf2.predict(X_test)))

[?] accuracy_score (train): 0.841189741661468
accuracy_score (test): 0.7975992263287505
```

```
print('precision_score (train):', precision_score(y_train, clf2.predict(X_train), average='macro'))
print('precision_score (test):', precision_score(y_test, clf2.predict(X_test), average='macro'))

[?] precision_score (train): 0.841189741661468
precision_score (test): 0.7975992263287505
```

```
print('recall_score (train):', recall_score(y_train, clf2.predict(X_train), average='micro'))
print('recall_score (test):', recall_score(y_test, clf2.predict(X_test), average='micro'))

[?] recall_score (train): 0.841189741661468
recall_score (test): 0.7975992263287505
```

Таким образом, на основании полученных с помощью метрик accuracy, precision_score и recall_score классификаций, можно сделать вывод о том, что для выбранного набора данных классификатор выполняет более качественную классификацию данных.

